

Wilson Vuu . Jarvis Consulting

As a Databricks Certified Data Engineer Associate with a background in software engineering, I'm currently at Jarvis, where I've been building Java backend applications and Python data pipelines end to end on Linux using PostgreSQL, SQL, and Azure Databricks, as well as containerizing and deploying applications in the cloud. I've been able to build on my previous experience at Cardata and Lowe-Martin, where I worked with backend APIs, databases, and automation using technologies like Node.js, Laravel, Python, and VBA to support financial reimbursement and logistics workflows. What draws me to software and data engineering is that there's always something new to learn or a better way to build something. Whether it's learning a new technology or improving an existing solution, I enjoy continuously learning and finding ways to make things better.

Skills

Proficient: Python, SQL/RDBMS, ETL Pipelines, Databricks, Spark/Pyspark, Java, Git

Competent: Azure, Docker, React.js, Linux/Bash, Node.js

Familiar: PHP (Laravel), Agile/Scrum, HTML/CSS, REST APIs, Javascript/TypeScript

Jarvis Projects

Project source code: https://github.com/jarviscanada/jarvis_data_eng_WilsonVuu

Cluster Monitor [GitHub]: The Linux Cluster Monitoring project is a system designed to track hardware specifications and resource usage across a networked Linux cluster. Each node collects its hardware profile on initialization and continuously logs CPU, memory, and disk metrics at scheduled intervals using crontab, with all data stored in a centralized PostgreSQL database managed through Docker. The data is organized into two tables covering hardware specifications and ongoing resource usage, giving the LCA team at Jarvis the visibility to monitor system performance over time and make informed decisions around resource allocation.

Kubernetes Cluster Monitor [GitHub]: Migrated the Linux Cluster Monitor to Azure, containerizing the monitoring agent with Docker and deploying it to Azure Kubernetes Service (AKS). Deployed the agent as a Kubernetes DaemonSet, so every node reports CPU, memory, and disk usage automatically, including new nodes as the cluster scales. Stored metrics in Azure Database for PostgreSQL and automated builds with Jenkins CI/CD, so changes are validated before reaching the production cluster.

Java Grep App [GitHub]: Implemented a Java application that replicates the Unix grep command, taking a regex pattern, a root directory, and an output file as arguments. The app recursively searches all files in the directory and writes matching lines to the output. Built using Java's Streams API with a lazy pipeline using Java NIO's Files.walk() for file traversal, flatMap to flatten each file's lines into a single stream, and filter to match the regex, ensuring memory stays low regardless of file or directory size. Used Maven for dependency and build management, packaging the app into a fat jar with all dependencies bundled. Containerized the application with Docker and pushed the image to DockerHub. Verified behavior with unit tests across edge cases.

Python Data Analytics [GitHub]: Analyzed two years of transactional data (2009-2011) for London Gift Shop (LGS), a UK-based wholesale retailer, to find actionable insights around customer retention and revenue growth. Applied RFM segmentation to identify high-value customers, quantify at-risk revenue, and support targeted win-back campaigns. Explored monthly sales trends, new vs. existing customer revenue splits, and customer health status over time to inform marketing and inventory decisions. Built using Python (pandas, NumPy, matplotlib, seaborn) in a Jupyter Notebook environment, with data stored in a PostgreSQL data warehouse, both running in Docker containers connected via a shared bridge network.

Databricks Data Engineering Project [GitHub]: Designed and implemented end-to-end data pipelines in Azure Databricks using PySpark and the Medallion Architecture (Bronze, Silver, Gold) to process transactional, customer, and stock market data from multiple sources. Built scalable ETL and DLT pipelines to ingest, cleanse, transform, and model data into Delta tables optimized for analytics. Delivered datasets and interactive dashboards that identified fraud patterns, customer behavior trends, and stock performance metrics.

Highlighted Projects

Inventory Management System: The Inventory Management System is a Spring Boot application with a React frontend, designed to help businesses organize their catalog and track stock levels. It models the relationships between

items, categories, and suppliers through a relational schema using Spring Data JPA/Hibernate, and exposes a RESTful API to create, update, and query inventory data across the catalog, delivering a clear, maintainable foundation for keeping stock accurate and operations running smoothly.

Smart Investors: Smart Investors is a web application with a JavaScript frontend that lets users practice investment strategies in a risk-free environment. It tracks each user's portfolio and transaction history through a relational database, and integrates the Yahoo Finance API to feed live market data into the platform for accurate portfolio valuations.

Professional Experiences

Software Developer/Data Engineer, Jarvis (2026-present): Developed and delivered a range of software and data engineering projects including a Linux cluster monitoring system, a Java grep application, and performed data analytics using Python in an Agile environment. Built backend services and data pipelines across Java and Python, working in Linux environments with tools such as Docker, PostgreSQL, and Bash.

Full Stack Software Engineer, Cardata (Nov 2021-Jul 2022): Integrated React.js and TypeScript frontends with Laravel and Node.js backend services to improve system scalability and the user experience of daily payment processing. Engineered a dynamic PDF export system to automate report generation and introduced a localization pipeline to the application, achieving 93% translation coverage.

Data/Programmer Analyst Intern, Lowe-Martin (May 2019-Sep 2019): Used VBA and Python to automate the sanitization of CSV data, resulting in an increase in shipping accuracy. Built custom report generation tools in Microsoft Access and Excel to display and present shipping and address data for internal review and validation. Built a prototype of the corporate website using HTML/CSS and JavaScript as a visual mockup for leadership to review and approve.

Education

Wilfrid Laurier University (2017-2022), Bachelor of Science, Computer Science

Miscellaneous

- Databricks Certified Data Engineer Associate (July 2026)
- Badminton player: Still a beginner, but I play in weekly recreational sessions. It's a fun way to stay active, and I really enjoy the quick pace and the social side of the game.
- Hiking enthusiast: I love hitting local trails and getting lost in conservation areas for a few hours. For me, it's a good way to clear my mind and recharge.